

Exp No: 1

INSTALLATION OF HADOOP

SCREENSHOTS

```

Apache Hadoop Distribution x + v
2024-08-13 22:24:48,714 INFO handler.ContextHandler: Stopped o.e.j.s.ServletContextHandler@47605f2f{/logs,/logs,file:///D:/Softwares/Hadoop/hadoop-3.4.0.tar/hadoop-3.4.0/logs/,STOPPED}
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: Stopping NameNode metrics system...
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: NameNode metrics system stopped.
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: NameNode metrics system shutdown complete.
2024-08-13 22:24:48,727 ERROR namenode.NameNode: Failed to start namenode.
org.apache.hadoop.hdfs.server.common.InconsistentFSStateException: Directory D:\hadoop-3.3.1\data\namenode is in an inconsistent state: storage directory does not exist or is not accessible.
    at org.apache.hadoop.hdfs.server.namenode.FSImage.recoverStorageDirs(FSImage.java:392)
    at org.apache.hadoop.hdfs.server.namenode.FSImage.recoverTransitionRead(FSImage.java:243)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.loadFSImage(FSNamesystem.java:1275)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.loadFromDisk(FSNamesystem.java:834)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.loadNamesystem(NameNode.java:817)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.initialize(NameNode.java:904)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:1158)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:1133)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.createNameNode(NameNode.java:1906)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.main(NameNode.java:1971)
2024-08-13 22:24:48,730 INFO util.ExitUtil: Exiting with status 1: org.apache.hadoop.hdfs.server.common.InconsistentFSStateException: Directory D:\hadoop-3.3.1\data\namenode is in an inconsistent state: storage directory does not exist or is not accessible.
2024-08-13 22:24:48,730 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at DESKTOP-SBBT9TG/192.168.56.1
*****/

D:\Softwares\Hadoop\hadoop-3.4.0.tar\hadoop-3.4.0\sbin>start-yarn.cmd
starting yarn daemons

D:\Softwares\Hadoop\hadoop-3.4.0.tar\hadoop-3.4.0\sbin>

```

```

Apache Hadoop Distribution x + v
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:54,318 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 8 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:57,344 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 9 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:59,367 WARN datanode.DataNode: Problem connecting to server: localhost/127.0.0.1:9000
2024-08-13 22:26:07,405 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 0 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:10,440 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 1 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:13,487 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 2 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:16,523 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 3 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:19,556 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 4 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:22,603 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 5 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:25,639 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 6 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:28,675 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 7 time(s);
2024-08-13 22:26:31,711 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 8 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:34,739 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 9 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:36,757 WARN datanode.DataNode: Problem connecting to server: localhost/127.0.0.1:9000
2024-08-13 22:26:44,808 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 0 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)

```

```
Apache Hadoop Distribution · X + v
2024-08-13 22:25:34,931 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-08-13 22:25:34,974 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class org.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false, ipcFailOver: false.
2024-08-13 22:25:34,980 INFO ipc.Server: Listener at 0.0.0.0:8030
2024-08-13 22:25:34,980 INFO ipc.Server: Starting Socket Reader #1 for port 8030
2024-08-13 22:25:34,989 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.api.ApplicationMasterProtocolPB to the server
2024-08-13 22:25:34,989 INFO ipc.Server: IPC Server Responder: starting
2024-08-13 22:25:34,989 INFO ipc.Server: IPC Server Listener on 8030: starting
2024-08-13 22:25:35,134 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class org.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false, ipcFailOver: false.
2024-08-13 22:25:35,134 INFO ipc.Server: Listener at 0.0.0.0:8032
2024-08-13 22:25:35,134 INFO ipc.Server: Starting Socket Reader #1 for port 8032
2024-08-13 22:25:35,147 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.api.ApplicationClientProtocolPB to the server
2024-08-13 22:25:35,147 INFO ipc.Server: IPC Server Responder: starting
2024-08-13 22:25:35,147 INFO ipc.Server: IPC Server Listener on 8032: starting
2024-08-13 22:25:36,299 INFO webproxy.ProxyCA: Created Certificate for OU=YARN-3fd21288-c809-4482-b2b1-c918b2ae3433
2024-08-13 22:25:36,380 INFO recovery.RMStateStore: Storing CA Certificate and Private Key
2024-08-13 22:25:36,380 INFO resourcemanager.ResourceManager: Transitioned to active state
2024-08-13 22:25:51,212 INFO resourcemanager.ResourceTrackerService: NodeManager from node DESKTOP-SBBT9TG(cmPort: 52991, httpPort: 8042) registered with capability: <memory:8192, vCores:8>, assigned nodeId DESKTOP-SBBT9TG:52991
2024-08-13 22:25:51,222 INFO rmnode.RMNodeImpl: DESKTOP-SBBT9TG:52991 Node Transitioned from NEW to RUNNING
2024-08-13 22:25:51,244 INFO capacity.AbstractLeafQueue: LeafQueue: root.default update max app related, maxApplications=10000, maxApplicationsPerUser=10000, Abs Cap:1.0, Cap: 1.0, MaxCap: 1.0
2024-08-13 22:25:51,244 INFO capacity.CapacityScheduler: Added node DESKTOP-SBBT9TG:52991 clusterResource: <memory:8192, vCores:8>
2024-08-13 22:25:51,244 INFO capacity.AbstractLeafQueue: LeafQueue: root.default update max app related, maxApplications=10000, maxApplicationsPerUser=10000, Abs Cap:1.0, Cap: 1.0, MaxCap: 1.0
```

```
Apache Hadoop Distribution · X + v
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver as a provider class
Aug 13, 2024 10:25:44 PM com.sun.jersey.server.impl.application.WebApplicationImpl _initiate
INFO: Initiating Jersey application, version 'Jersey: 1.19.4 05/24/2017 03:20 PM'
Aug 13, 2024 10:25:44 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver to GuiceManagedComponentProvider with the scope "Singleton"
Aug 13, 2024 10:25:45 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope "Singleton"
Aug 13, 2024 10:25:45 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.NMWebServices to GuiceManagedComponentProvider with the scope "Singleton"
2024-08-13 22:25:45,534 INFO handler.ContextHandler: Started o.e.j.w.WebAppContext@a1bb14{node/, file:///C:/Users/Hares/h/AppData/Local/Temp/jetty-0_0_0-8042-hadoop-yarn-common-3_4_0-jar-any-3589810351802953548/webapp/,AVAILABLE!}{jar:file:/D:/Softwares/Hadoop/hadoop-3.4.0.tar/hadoop-3.4.0/share/hadoop/yarn/hadoop-yarn-common-3.4.0.jar!/webapps/node}
2024-08-13 22:25:45,558 INFO server.AbstractConnector: Started ServerConnector@4cdb8504{HTTP/1.1, (http/1.1)}{0.0.0.0:8042}
2024-08-13 22:25:45,558 INFO server.Server: Started @28007ms
2024-08-13 22:25:45,560 INFO webapp.WebApps: Web app node started at 8042
2024-08-13 22:25:45,560 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : DESKTOP-SBBT9TG:52991.
2024-08-13 22:25:45,565 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-08-13 22:25:45,723 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8031
2024-08-13 22:25:45,806 INFO nodemanager.NodeStatusUpdaterImpl: Running Applications Size : 0.
2024-08-13 22:25:51,244 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id -165124043
2024-08-13 22:25:51,244 INFO security.NMTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id 88793460
2024-08-13 22:25:51,244 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as DESKTOP-SBBT9TG:52991 with total resource of <memory:8192, vCores:8>
```

hadoop

Cluster

About Nodes Node Labels Annotations

NEW NEW SAVING SUBMITTED ACCEPTED RUNNING FINISHED FAILED KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Used Resources	Total Resources	Reser
0	0	0	0	0	<memory:0 B, vCores:0>	<memory:8 GB, vCores:8>	<memory:0 B, vCores:0>

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes
1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority	Scheduler Busy %
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0	0

Show 20 entries

ID	User	Name	Application Type	Application Tags	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU Vcores	Allocated Memory MB	Allocated GPUs	Reserved CPU Vcores
No data available in table																

Showing 0 to 0 of 0 entries